

# Segmenting the Right-of-Way of San Francisco

Using VGG16-UNET on RGB Satellite Imagery

MUSA 650 Final Project - Alexander Helms

## Methods

Since I wanted to use image segmentation on a large RGB dataset, I needed to pick a model and architecture that will:

- 2 be generalized enough to not overfit on
  - *street-fog* shapes
  - chaotic assortments of cars, trees, & such
  - potentially unclassified or incorrect RoW
- 3 not be too process intensive

The model and architecture I ended up using is an odd VGG16 and UNet hybrid that I found from [a journal by Pravatiasari et al.](#) and in a [related article by Gazali](#). Both are using this architecture to perform segmentation of MRI scans to find brain tumors. Although the use case is different, the concept and model are similar.

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')

import os
import pandas as pd
import json
import numpy as np
import dataframes as dd

PROJ_DIR = 'drive/MyDrive/Penn/MUSA-650/FinalProject'
tod [PROJ_DIR]

#DATA DIR = PROJ_DIR + '/' + 'data'
DATA_DIR = 'data'

IMAGE_PATH = DATA_DIR + '/clean/' + 'sf_naip_images.parquet'
MASK_PATH = DATA_DIR + '/clean/' + 'sf_naip_mask.parquet'
INFO_PATH = DATA_DIR + '/clean/' + 'sf_naip_info.parquet'

!python -m pip install 'fsspec>0.3.3'

# CLEAN DATAFRAMES AND TURN THEM TO ARRAYS
def df_to_array(arr_df):
    # import csv of image data
    # get file names
    cols = [col for col in list(arr_df) if 'Unnamed' not in col]
    arr_df = arr_df[cols]

    # list of shape index colnames
    icols = [col for col in list(arr_df) if 'aes' not in col]
    # make array out of rgb dataframe
    #img_arr = np.stack(arr_df[icols]).apply(list, axis=1)
    #img_arr = arr_df[icols].to_numpy(dtype='int')
    print('Shape: {}'.format(img_arr.shape))

    return img_arr

# UNFLATTEN FUNCTION
def unflatten(array, new_shape):
    # shape initial
    n_samples = array.shape[0]
    flat_shape = array.shape[1]

    new_shape = [n_samples] + list(new_shape)

    return array.reshape(new_shape)

# SHAPES
img_shape = (128, 128, 4)
mask_shape = (128, 128, 1)
rgb_shape = (128,128,3)

# IMPORT IMAGE ARRAYS
img_dd = dd.read_parquet(
    IMAGE_PATH,
    blocksize=1000000,
    sample=55350,
    dtype='int'
)
img_arr = df_to_array(img_dd.compute(scheduler='processes'))

# IMPORT ROW MASKS
mask_arr = df_to_array(pd.read_parquet(MASK_PATH))

# IMPORT ADDITIONAL INFO
info_df = pd.read_parquet(INFO_PATH)
```

## Train / Test Split

When preparing the data, I perform a 75 / 25 split on the data. I can afford 25% for training as I have a large dataset.

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential, Model
from tensorflow.keras.utils import plot_model
from tensorflow.keras.optimizers import Adam, RMSprop
import keras

# load vgg model
from keras.applications.vgg16 import VGG16, preprocess_input

# SPLIT GREY DATAFRAME INTO 50/50
WITH STRATIFICATION ON THE CLASSES COL
X_train, X_test, y_train, y_test = train_test_split(
    X_train, mask_arr,
    img_arr, mask_arr,
    test_size=0.75, random_state=420
)

## Scale the data
scaler = MinMaxScaler()
scaler.fit(X_train)
#X_train = scaler.transform(X_train)
#X_test = scaler.transform(X_test)
X_train = unflatten(X_train, img_shape)[:,:,:,:3]
#X_train = reduce_sum(y_train)
X_test = unflatten(X_test, img_shape)[:,:,:,:3]
#X_test = preprocess_input(X_test)
y_train = unflatten(y_train, mask_shape)
y_test = unflatten(y_test, mask_shape)
print('X train: {}'.format(X_train.shape))
print('Y train: {}'.format(y_train.shape))

X_train: (4702, 128, 128, 3)
Y_train: (4702, 128, 128, 1)
```

## Model

As mentioned, I am basing my architecture off of an article by Gazali [\[source\]](#). I found this architecture very impactful as U-Net provides a helpful path for semantic segmentation. By narrowing the shape of the data then bringing it back out to the same shape, it is able to pool the connections into buckets. When the image is resized back to its original shape, it is then able to take with the classification for each pixel!

## Import VGG16 Model

```
In [ ]: # load the model
vgg16 = VGG16(
    include_top = False,
    input_shape = rgb_shape,
    weights='imagenet'
)

In [ ]: import tensorflow as tf
from tensorflow.keras.layers import \
    Conv2D, Conv2DTranspose, MaxPooling2D, \
    Dense, Activation, Dropout, \
    Flatten, Input, Concatenate, \
    BatchNormalization
from tensorflow.keras.initializers import RandomNormal

def conv_block(inputs, num_filters):
    x = Conv2D(num_filters, 3, padding='same')(inputs)
    # BatchNormalization(input=y_train)
    x = Activation('relu')(x)
    x = Conv2D(num_filters, 3, padding='same')(x)
    x = BatchNormalization(x)
    x = Activation('relu')(x)
    return x

def define_decoder(inputs, skip_layer, num_filters):
    init = RandomNormal(stddev=0.02)
    x = Conv2DTranspose(
        num_filters, (2,2),
        strides=(2,2),
        padding='same',
        kernel_initializer=init
    )(inputs)
    g = Concatenate()([x, skip_layer])
    g = conv_block(g, num_filters)
    return g

def vgg16_unet(input_shape):
    inputs = Input(shape=input_shape)
    vgg16 = VGG16(include_top=False, weights='imagenet', input_tensor=inputs)
    # No need to extract encoder based on their output shape from vgg16 model
    s1 = vgg16.get_layer('block1_conv2').output
    s2 = vgg16.get_layer('block2_conv2').output
    s3 = vgg16.get_layer('block3_conv3').output
    s4 = vgg16.get_layer('block4_conv3').output
    b1 = vgg16.get_layer('block5_conv3').output # bottleneck/bridge layer from vgg16
    b2 = vgg16.get_layer('block5_conv3').output #32

    # Decoder Block
    d1 = define_decoder(d1,s4,s12)
    d2 = define_decoder(d1,s3,s256)
    d3 = define_decoder(d2,s2,s128)
    d4 = define_decoder(d3,s1,s64) #output layer
    outputs = Conv2D(1,1, padding='same', activation='sigmoid')(d4)
    model = Model(inputs, outputs)

    return model
```

## Import Metrics

```
In [ ]: from tensorflow.math import reduce_sum

# Dice Loss = (2 * intersection) / (true_positive + false_positive + false_negative)
# a measure of overlap between two samples
# ranges from 0 to 1
# a dice coefficient of 1 denotes perfect and complete overlap
def dice_coef(y_true, y_pred):
    y_true = Flatten()(y_true)
    y_pred = Flatten()(y_pred)
    intersection = reduce_sum(y_true*y_pred)
    return (2. * intersection + smooth) / (reduce_sum(y_true) + reduce_sum(y_pred))

def dice_loss(y_true, y_pred):
    return 1.0 - dice_coef(y_true, y_pred)

# Intersection-Over-Union
def true_positive / (true_positive + false_positive + false_negative)
def iou(y_true, y_pred):
    intersection = (y_true*y_pred).sum()
    union = y_true.sum() + y_pred.sum() - intersection
    x = (intersection + 1e-15) / (union + 1e-15)
    x = x.astype(np.float32)
    return x
return tf.numpy_function(tf, [y_true, y_pred], tf.float32)
```

## Compile Model

```
In [ ]: from tensorflow.keras.metrics import Precision, Recall

model = vgg16_unet(rgb_shape)
model.compile(
    loss=dice_loss,
    optimizer=Adam(lr = .001),
    metrics=[dice_coef, iou, Recall(), Precision()])

/usr/local/lib/python3.7/dist-packages/keras/optimizer_v2/adam.py:105: UserWarning: The 'lr' argument is deprecated, use 'learning_rate' instead.
  super(Adam, self)._init(name, **kwargs)
```

```
In [45]: model.summary()

Model: "model"
Layer (type) Output Shape Param # Connected to
-----
input_2 (InputLayer) [(None, 128, 128, 3) 0] []
block1_conv1 (Conv2D) (None, 128, 128, 64) 1792 ['input_2[0][0][0]']
block1_conv2 (Conv2D) (None, 128, 128, 64) 36928 ['block1_conv1[0][0][0]']
block1_pool (MaxPooling2D) (None, 64, 64, 128) 0 ['block1_conv2[0][0][0]']
block2_conv1 (Conv2D) (None, 64, 64, 128) 73856 ['block1_pool[0][0][0]']
block2_conv2 (Conv2D) (None, 64, 64, 128) 147584 ['block2_conv1[0][0][0]']
block2_pool (MaxPooling2D) (None, 32, 32, 128) 0 ['block2_conv2[0][0][0]']
block3_conv1 (Conv2D) (None, 32, 32, 256) 295168 ['block2_pool[0][0][0]']
block3_conv2 (Conv2D) (None, 32, 32, 256) 590080 ['block3_conv1[0][0][0]']
block3_conv3 (Conv2D) (None, 32, 32, 256) 590080 ['block3_conv2[0][0][0]']
block3_pool (MaxPooling2D) (None, 16, 16, 256) 0 ['block3_conv3[0][0][0]']
block4_conv1 (Conv2D) (None, 16, 16, 512) 1180160 ['block3_pool[0][0][0]']
block4_conv2 (Conv2D) (None, 16, 16, 512) 2359808 ['block4_conv1[0][0][0]']
block4_conv3 (Conv2D) (None, 16, 16, 512) 2359808 ['block4_conv2[0][0][0]']
block4_pool (MaxPooling2D) (None, 8, 8, 512) 0 ['block4_conv3[0][0][0]']
block5_conv1 (Conv2D) (None, 8, 8, 512) 2359808 ['block5_conv1[0][0][0]']
block5_conv2 (Conv2D) (None, 8, 8, 512) 2359808 ['block5_conv2[0][0][0]']
conv2d_transpose (Conv2DTransp (None, 16, 16, 512) 1049088 ['block5_conv3[0][0][0]']
concatenate (Concatenate) (None, 16, 16, 1024) 0 ['conv2d_transpose[0][0][0]', 'block1_conv2[0][0][0]']
conv2d (Conv2D) (None, 16, 16, 512) 4719104 ['concatenate[0][0][0]']
batch_normalization (BatchNorm (None, 16, 16, 512) 2048 ['conv2d[0][0][0]']
activation (Activation) (None, 16, 16, 512) 0 ['batch_normalization[0][0][0]']
conv2d_1 (Conv2D) (None, 16, 16, 512) 2359808 ['activation[0][0][0]']
batch_normalization_1 (BatchNo (None, 16, 16, 512) 2048 ['conv2d_1[0][0][0]']
activation_1 (Activation) (None, 16, 16, 512) 0 ['batch_normalization_1[0][0][0]']
conv2d_transpose_1 (Conv2DTran (None, 32, 32, 256) 524544 ['activation_1[0][0][0]']
concatenate_1 (Concatenate) (None, 32, 32, 512) 0 ['conv2d_transpose_1[0][0][0]', 'block1_conv2[0][0][0]']
conv2d_2 (Conv2D) (None, 32, 32, 256) 1179904 ['concatenate_1[0][0][0]']
batch_normalization_2 (BatchNo (None, 32, 32, 256) 1024 ['conv2d_2[0][0][0]']
activation_2 (Activation) (None, 32, 32, 256) 0 ['batch_normalization_2[0][0][0]']
conv2d_3 (Conv2D) (None, 32, 32, 256) 590080 ['activation_2[0][0][0]']
batch_normalization_3 (BatchNo (None, 32, 32, 256) 1024 ['conv2d_3[0][0][0]']
activation_3 (Activation) (None, 32, 32, 256) 0 ['batch_normalization_3[0][0][0]']
conv2d_transpose_2 (Conv2DTran (None, 64, 64, 128) 131200 ['activation_3[0][0][0]']
concatenate_2 (Concatenate) (None, 64, 64, 256) 0 ['conv2d_transpose_2[0][0][0]', 'block1_conv2[0][0][0]']
conv2d_4 (Conv2D) (None, 64, 64, 128) 295040 ['concatenate_2[0][0][0]']
batch_normalization_4 (BatchNo (None, 64, 64, 128) 512 ['conv2d_4[0][0][0]']
activation_4 (Activation) (None, 64, 64, 128) 0 ['batch_normalization_4[0][0][0]']
conv2d_5 (Conv2D) (None, 64, 64, 128) 147584 ['activation_4[0][0][0]']
batch_normalization_5 (BatchNo (None, 64, 64, 128) 512 ['conv2d_5[0][0][0]']
activation_5 (Activation) (None, 64, 64, 128) 0 ['batch_normalization_5[0][0][0]']
conv2d_transpose_3 (Conv2DTran (None, 128, 128, 64) 32832 ['activation_5[0][0][0]']
concatenate_3 (Concatenate) (None, 128, 128, 128) 0 ['conv2d_transpose_3[0][0][0]', 'block1_conv2[0][0][0]']
conv2d_6 (Conv2D) (None, 128, 128, 64) 73792 ['concatenate_3[0][0][0]']
batch_normalization_6 (BatchNo (None, 128, 128, 64) 256 ['conv2d_6[0][0][0]']
activation_6 (Activation) (None, 128, 128, 64) 0 ['batch_normalization_6[0][0][0]']
conv2d_7 (Conv2D) (None, 128, 128, 64) 36928 ['activation_6[0][0][0]']
batch_normalization_7 (BatchNo (None, 128, 128, 64) 256 ['conv2d_7[0][0][0]']
activation_7 (Activation) (None, 128, 128, 64) 0 ['batch_normalization_7[0][0][0]']
conv2d_8 (Conv2D) (None, 128, 128, 1) 65 ['activation_7[0][0][0]']
Total params: 25,862,337
Trainable params: 25,858,497
Non-trainable params: 3,840
```

## Fit Model

```
In [ ]: tf.config.run_functions_eagerly(True)

In [ ]: EPOCHS = 25
#train_steps = len(X_train)//batch_size
#test_steps = len(X_test)//batch_size

history = model.fit(
    X_train, tf.cast(y_train, tf.float32),
    epochs=EPOCHS,
    validation_data=(X_test, tf.cast(y_test, tf.float32)),
    callbacks=[tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', patience=5, verbose=1),
               tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=10, verbose=1),
               tf.keras.callbacks.TensorBoard(log_dir='./logs')],
    verbose=1)

# Save the history.history dict to a pandas DataFrame:

/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/structured_function.py:265: UserWarning: Even though the 'tf.config.experimental_run_functions_eagerly' option is set, this option does not apply to tf.data functions. To force eager execution of tf.data functions, please use 'tf.data.experimental.enable_debug_mode()'
  "Even though the 'tf.config.experimental_run_functions_eagerly' "
147/147 [=====] - 4432s 30s/step - loss: 0.2445 - dice_coef: 0.7555 - iou: 0.6095 - re
call: 0.7806 - precision: 0.7519 - val_loss: 0.2305 - val_dice_coef: 0.7695 - val_iou: 0.6267 - val_recall: 0.7
710 - val_precision: 0.7932
Epoch 2/25
147/147 [=====] - 4356s 30s/step - loss: 0.1956 - dice_coef: 0.8044 - iou: 0.6744 - re
call: 0.8102 - precision: 0.8070 - val_loss: 0.2085 - val_dice_coef: 0.7915 - val_iou: 0.6565 - val_recall: 0.7
1710 - val_precision: 0.8895
Epoch 3/25
147/147 [=====] - 4371s 30s/step - loss: 0.1679 - dice_coef: 0.8821 - iou: 0.7138 - re
call: 0.8293 - precision: 0.8399 - val_loss: 0.2650 - val_dice_coef: 0.7350 - val_iou: 0.5825 - val_recall: 0.7
972 - val_precision: 0.6862
Epoch 4/25
147/147 [=====] - 4364s 30s/step - loss: 0.1549 - dice_coef: 0.8452 - iou: 0.7333 - re
call: 0.8198 - precision: 0.8568 - val_loss: 0.1597 - val_dice_coef: 0.8413 - val_iou: 0.7213 - val_recall: 0.8
707 - val_precision: 0.8171
Epoch 5/25
147/147 [=====] - 4375s 30s/step - loss: 0.1415 - dice_coef: 0.8583 - iou: 0.7533 - re
call: 0.8518 - precision: 0.8653 - val_loss: 0.1509 - val_dice_coef: 0.8491 - val_iou: 0.7390 - val_recall: 0.8
249 - val_precision: 0.8779
Epoch 6/25
147/147 [=====] - 4359s 30s/step - loss: 0.1326 - dice_coef: 0.8674 - iou: 0.7668 - re
call: 0.8659 - precision: 0.8805 - val_loss: 0.1719 - val_dice_coef: 0.8281 - val_iou: 0.7085 - val_recall: 0.7
606 - val_precision: 0.9123
Epoch 7/25
147/147 [=====] - 4357s 30s/step - loss: 0.1257 - dice_coef: 0.8743 - iou: 0.7778 - re
call: 0.8650 - precision: 0.8859 - val_loss: 0.1368 - val_dice_coef: 0.8590 - val_iou: 0.7542 - val_recall: 0.8
116 - val_precision: 0.9149
Epoch 8/25
147/147 [=====] - 4358s 30s/step - loss: 0.1252 - dice_coef: 0.8747 - iou: 0.7787 - re
call: 0.8652 - precision: 0.8896 - val_loss: 0.1368 - val_dice_coef: 0.8632 - val_iou: 0.7606 - val_recall: 0.8
405 - val_precision: 0.8896
Epoch 9/25
58/147 [=====] - ETA: 26:17 - loss: 0.1217 - dice_coef: 0.8783 - iou: 0.7838 - recal
l: 0.8728 - precision: 0.8867
```

```
In [ ]: import datetime

now = datetime.datetime.now()
dt_string = now.strftime("%m%d_%H%M")

model_path = os.path.join(DATA_DIR, "models", 'vgg16_model_{}.keras'.format(dt_string))
model_save(model_path)

hist_df = pd.DataFrame(history.history)
history_path = os.path.join(DATA_DIR, "models", 'vgg16_history_{}.csv'.format(dt_string))
with open(history_path, mode='w') as f:
    hist_df.to_csv(f)
```

## Results

### Accuracy

Instead of accuracy, I used two different metrics: a Dice Coefficient and an Intersection over Union.

Dice Loss is able to measure the overlap between two samples. In my case, that's the observed Right-of-Way Mask/Training set and my predicted RoW. The coefficient closer to one means there is a perfect overlap.

Looking at the middle plot of the figure below, the testing Dice Coefficient (blue) was able to rise to .85 by the 5th epoch. Which is amazing. The Validation was more inconsistent, but eventually rose to .9. The inconsistency could mean that I don't have that generalized of a model. I could use a dropout layer to add more variety to my dataset.

Intersection-Over-Union is a common evaluation metric for semantic image segmentation. It also measures the overlap but uses bounding boxes instead.

It is helpful to look at both the Dice Loss and the IoU as they give different assessments. The Dice Loss primarily looks at the overlap, regardless of the shape. But might not be the best at getting a larger context of the shape. The IoU metric can help understand the larger positioning.

The IoU metric, based on the figure below, took longer to rise than the dice coefficient. This is likely due to a number of reasons. Missing RoW polygon masks could lower the metrics as the model will still predict a road. The IoU metric will be affected more by this missing mask as it compares more area.

Regardless, both metrics got fairly close to .9, which suggests that the model is fairly accurate.

```
In [39]: def plot_loss_acc(fitted_model, num_epochs = EPOCHS):
import matplotlib.pyplot as plt
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
ax[2].set_xlabel('epochs')
```

```
ax[0].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='blue', label='Dice Coefficient')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_dice_coef'], c='red', label='Validation Dice Coef')
ax[2].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(1, num_epochs+1), fitted_model.history['loss'], c='blue', label='Training loss')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_loss'], c='red', label='Validation loss')
ax[2].plot(range(1, num_epochs+1), fitted_model.history['dice_coef'], c='red', label='Validation Dice Coef')
ax[0].legend()
ax[1].set_xlabel('epochs')
```

```
ax[2].plot(range(1, num_epochs+1), fitted_model.history['iou'], c='blue', label='Intersection Over Union')
ax[1].plot(range(1, num_epochs+1), fitted_model.history['val_iou'], c='red', label='Validation IoU')
ax[2].legend()
ax[2].set_xlabel('epochs')
```

```
plot_loss_acc(model.history, EPOCHS)
```

```
fig, ax = plt.subplots(1,3, figsize=(24,6))
ax[0].plot(range(
```